

Recommender Systems class

Evaluation and testing



Evaluation measures

Regression

- MSE
- RMSE
- MAE
- MAPE (MRE)
- TRE

Classification

- Sensitivity/Recall/True Positive Rate
- Precision
- Accuracy
- F1 score

Ranking

- HR@n
- NDCG@n
- MAP@n

Regression evaluation measures

- MAE – Mean Absolute Error

$$\frac{1}{n} \sum_{i=1}^n |\hat{y}_i - y_i|$$

- MSE – Mean Squared Error

$$\frac{1}{n} \sum_{i=1}^n (\hat{y}_i - y_i)^2$$

- RMSE – Root Mean Squared Error

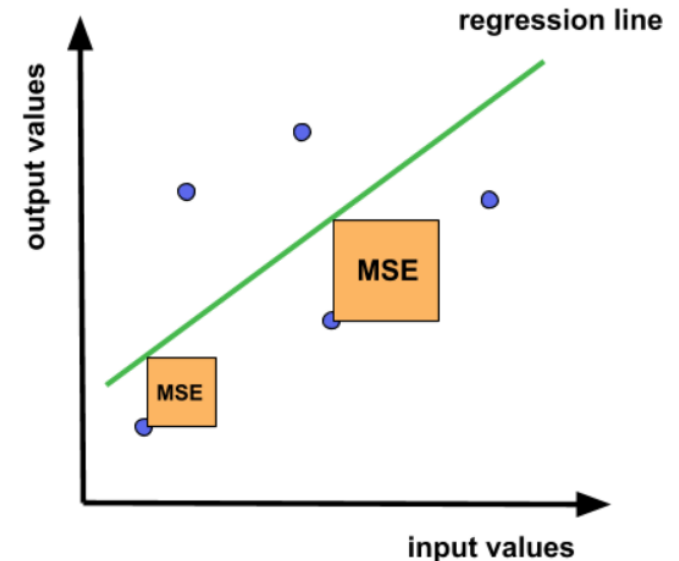
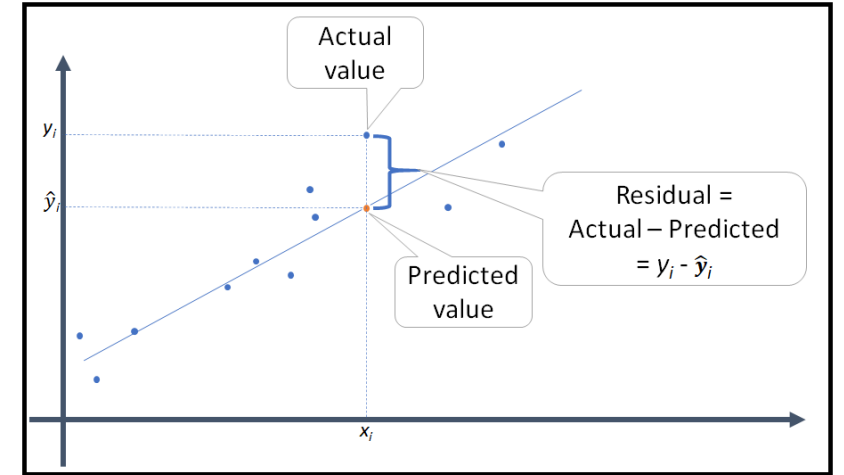
$$\sqrt{\frac{1}{n} \sum_{i=1}^n (\hat{y}_i - y_i)^2}$$

- MAPE (MRE) – Mean Absolute Percentage Error (Mean Relative Error)

$$\frac{1}{n} \sum_{i=1}^n \frac{|\hat{y}_i - y_i|}{|y_i|}$$

- TRE – Total Relative Error

$$\frac{\sum_{i=1}^n |\hat{y}_i - y_i|}{\sum_{i=1}^n |y_i|}$$



Classification evaluation measures

- Sensitivity/Recall/True Positive Rate

$$\frac{TP}{TP + FN}$$

- Precision

$$\frac{TP}{TP + FP}$$

- Accuracy

$$\frac{TP + TN}{P + N}$$

- F1 score


$$\frac{2}{\frac{1}{Sensitivity} + \frac{1}{Precision}}$$

Close to one – good

Close to zero – bad

		True condition	
		Condition positive	Condition negative
Predicted condition	Predicted condition positive	True positive	False positive, Type I error
	Predicted condition negative	False negative, Type II error	True negative

ROC curve



Ranking evaluation measures

- HR@n – Hit Ratio

$$\frac{1}{|U|} \sum_{u \in U} \sum_{k=1}^n 1_{D_u}(r_{u,k})$$

- NDCG@n – Normalized Discounted Cumulative Gain

$$\frac{1}{|U|} \sum_{u \in U} \sum_{k=1}^n \frac{1_{D_u}(r_{u,k})}{\log_2(1 + k)}$$

- MAP@n – Mean Average Precision

$$\frac{1}{|U|} \sum_{u \in U} \frac{1}{|D_u|} \sum_{k=1}^n \frac{\text{HR}@k}{k} 1_{D_u}(r_{u,k})$$

D_u - items user u actually interacted with
 $r_{u,k}$ - k -th recommendation for user u

Position	Movie	Score
1	Rocky	0.98
2	Interstellar	0.86
3	Shrek	0.83
4	Shawshank Redemption	0.75
5	Lion King	0.69
6	Star Wars	0.61
7	Apocalypto	0.55

$$\text{HR}@7 = 1 + 0 + 1 + 1 + 0 + 0 + 1 = 4$$

$$\text{NDCG}@7 = \frac{1}{\log_2 2} + \frac{0}{\log_2 3} + \frac{1}{\log_2 4} + \frac{1}{\log_2 5} + \frac{0}{\log_2 6} + \frac{0}{\log_2 7} + \frac{1}{\log_2 2} = 2.26401$$

$$\text{MAP}@7 = \left(\frac{1}{1} \cdot 1 + \frac{1}{2} \cdot 0 + \frac{2}{3} \cdot 1 + \frac{3}{4} \cdot 1 + \frac{3}{5} \cdot 0 + \frac{3}{6} \cdot 0 + \frac{4}{7} \cdot 1 \right) / 4 = 0.7470238095$$

Other evaluation measures

Coverage

- The percentage of all available items in the first n recommendations for all users

Novelty

- Evaluates the likelihood that the user was not aware of the recommended items

Serendipity

- Should measure the surprise effect in recommendations

Diversity

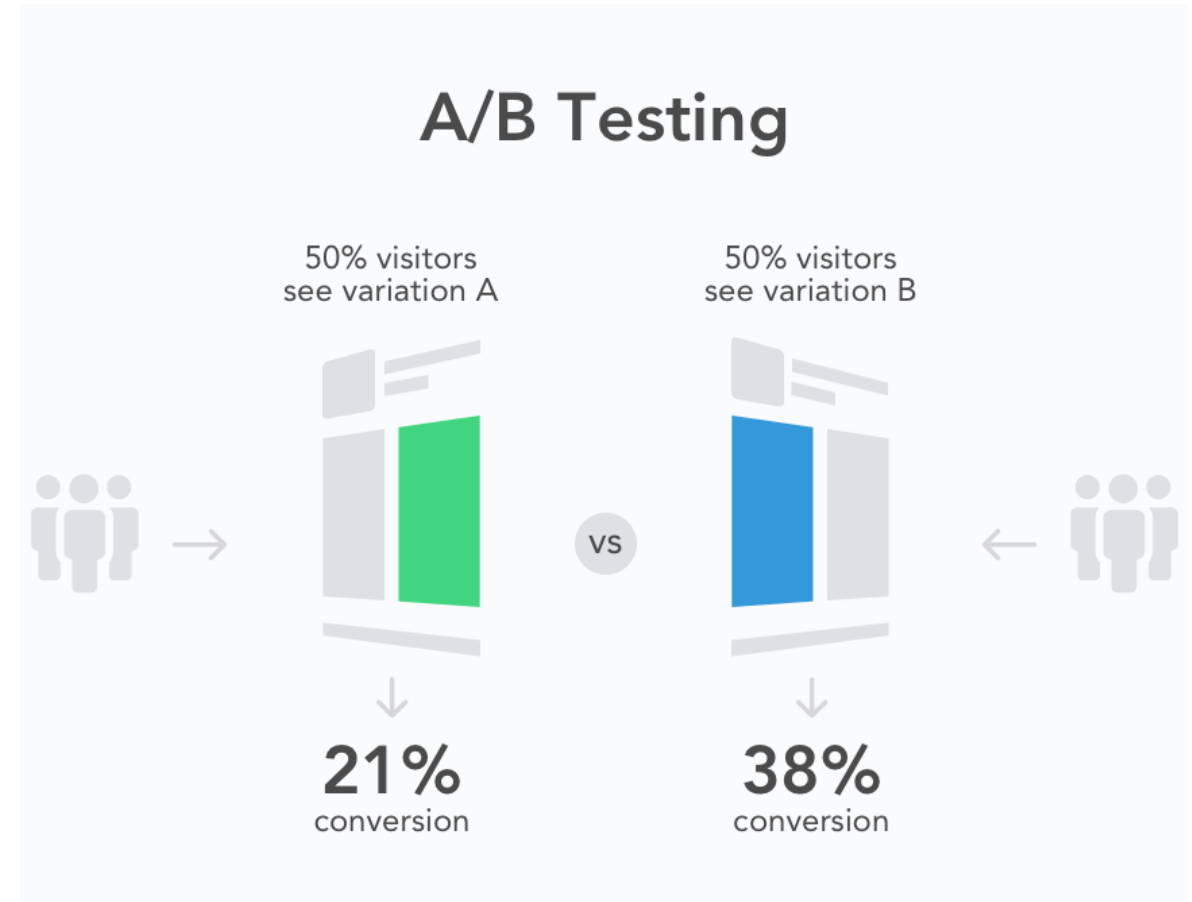
- The first n recommendations should be diverse enough so that if the user does not like the first item he/she might still like the other recommendations

Testing schemes

- Online
 - A/B tests
 - Counterfactuals
- Offline
 - Train-test split
 - Train-validation-test split
 - Cross-validation
 - Leave-one-out
 - Simulation

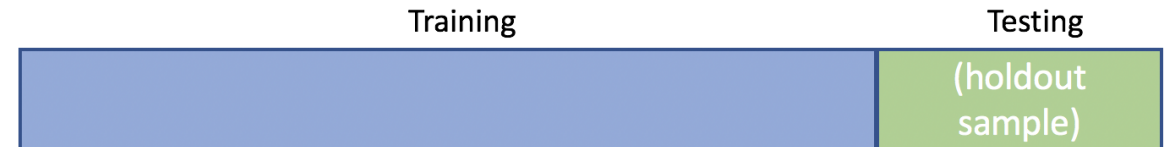
A/B tests

- Assumes that a model/algorithm is already used in production
- Split all cases when the algorithm is used into two groups: A and B (typically 50%-50%)
- After a predefined time gather and compare results of both models/algorithms
- Examples:
 - Two recommenders on YouTube – users are randomly split in half
 - Two trading algorithms – budget is split in half
 - Two website versions are presented to users – cookies are randomly assigned to both groups



Train-test split

- Divide the dataset into two parts:
 - training set
 - test set
- Train the model on the training set
- Evaluate the model on the test set



- + Good when the dataset is large and training is expensive
- A single dataset split may not properly reflect model's ability to generalize

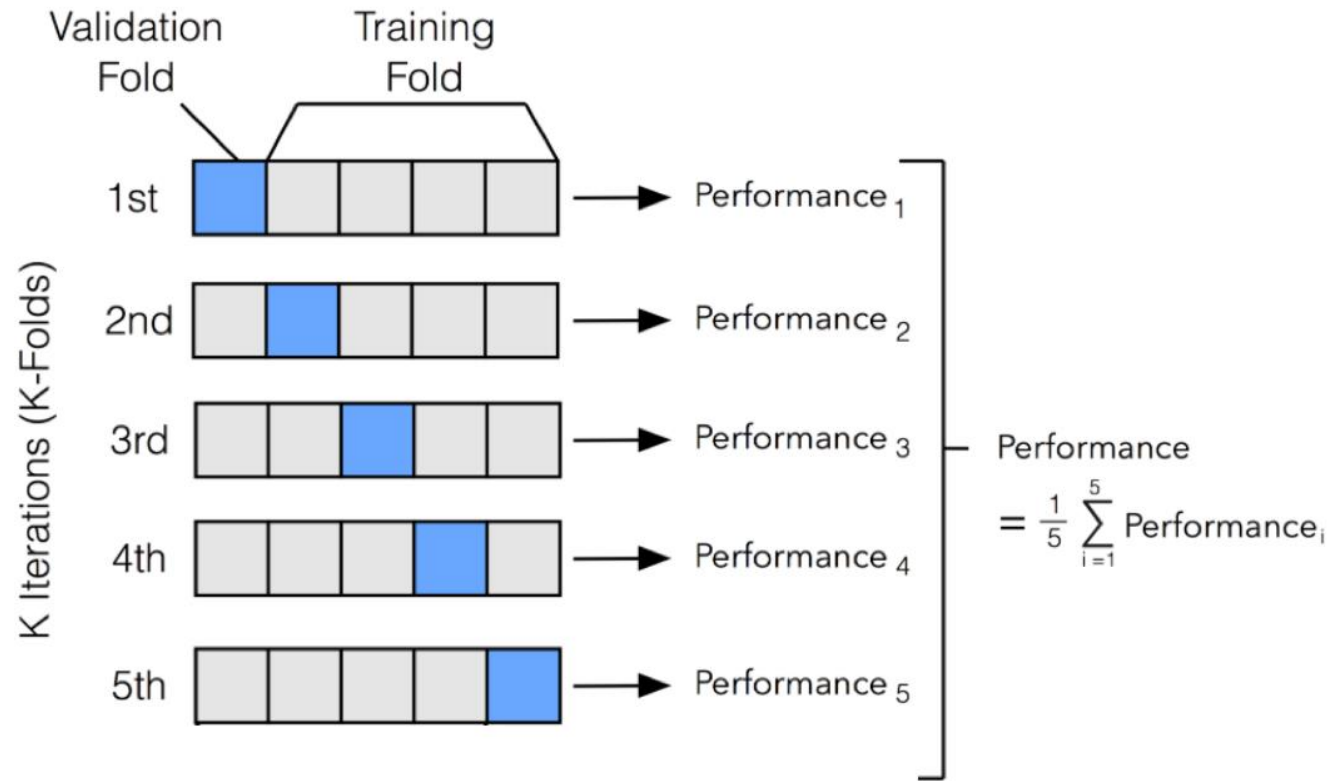
Train-validation-test split

- Divide the dataset into three parts:
 - training set
 - validation set
 - test set
 - Train the model on the training set with many sets of hyperparameters
 - Evaluate all trained models on the validation set
 - Choose hyperparameters which give the best result
 - Evaluate the final model on the test set to check for overfitting
-
- + Good when the dataset is large and training is expensive
 - + Allows for a proper hyperparameter tuning
 - A single dataset split may not properly reflect model's ability to generalize



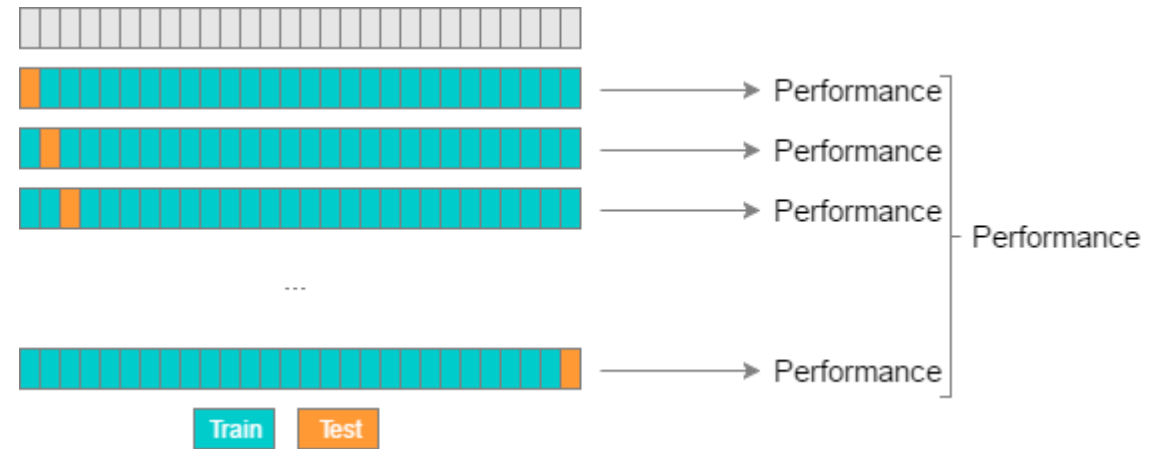
Cross-validation

- Divide the dataset D into K equal-sized parts
 - Train and evaluate the model K times in the following way:
 - Choose the i -th dataset part D_i
 - Train the model on $D \setminus D_i$
 - Evaluate the model on D_i
 - Gather all K results and aggregate them into a single measure (for instance by taking an average)
- + Allows to test the model on the entire dataset removing the risk of choosing an atypical split
- When K is large and model training is expensive, this method requires a long processing time



Leave-one-out

- Leave-one-out is an extreme case of cross-validation where K is equal to the number of elements in the dataset
- Train and evaluate the model K times in the following way:
 - Choose the i -th element $y_i \in D$
 - Train the model on $D \setminus \{y_i\}$
 - Evaluate the model on y_i
- Gather all results and aggregate them into a single measure



- + Allows to test the model on the entire dataset removing the risk of choosing an atypical split
- + Good when the training set is small
- + Good when the typical proportion of the production training dataset size is large compared to the number of predictions to be made
- For even moderately complex models requires a long processing time

Explicit feedback vs implicit feedback testing

Explicit feedback, e.g. ratings

- Treat the recommender as a typical regressor and use regression measures to evaluate it
- Generate prediction for every pair user-item in the test set
- You can simplify the testing scheme by generating a single prediction at a time
- When serving recommend items with the highest predictions first

Implicit feedback, e.g. binary indicators if there was an interaction or not

- Generate a set of recommendations for each user in the test set (typically by assigning a score to every item in the set of items the user have not interacted with)
- Use ranking evaluation measures to take positions into account